







The **Mixer** slot is where attention / CoTAR / **MI (ours)** are swapped. The MI variants (v1–v4) are detailed below.

MI v1 — fixed PAC coupling as aggregation weights (no learned Q/K/V)

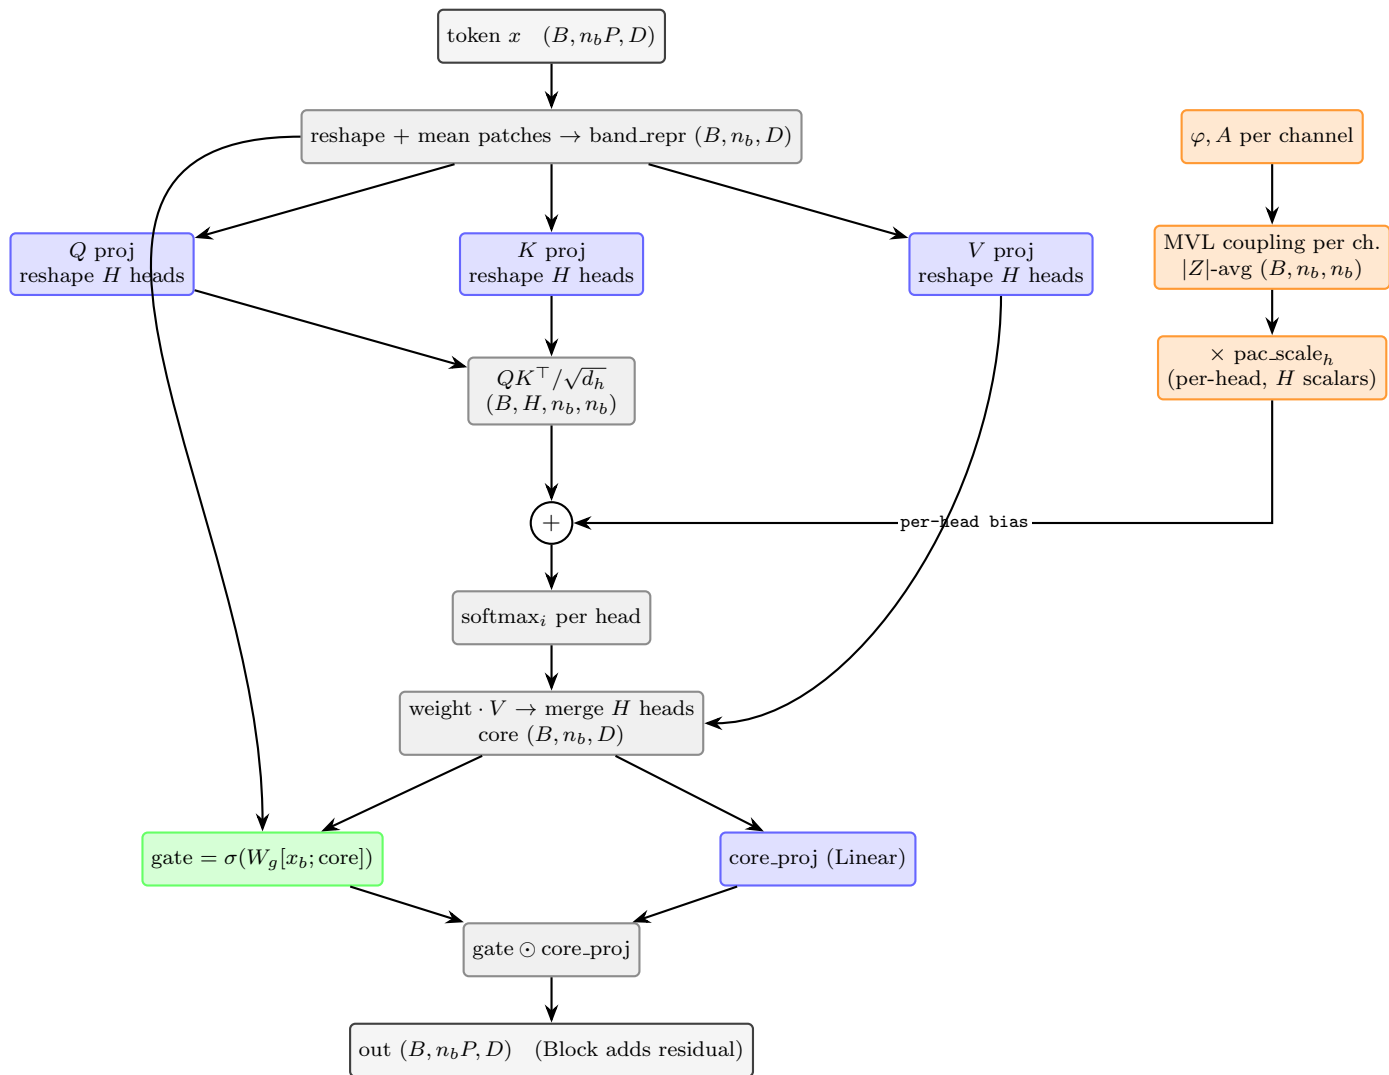




MI v3 — PAC-biased attention, *per-channel* coupling + seed-controlled (beats CoTAR 3/3 seeds)



MI v4 — multi-head + per-head pac_scale + gated injection (tried, no net gain, reverted)



MI v5 — adaptive gated PAC branch over a full self-attention floor (current; $g \rightarrow 0$ = attention baseline exactly)

floor
= vanilla
attention
baseline

